

Karp-Rabin fingerprinting on strings

Given a string $S = S[0 \dots n-1]$, and two positions $0 \leq i < j \leq n-1$, the longest common extension $lce_S(i, j)$ is the length of the maximal run of matching characters from those positions, namely: if $S[i] = S[j]$ then $lce_S(i, j) = 0$; otherwise, $lce_S(i, j) = \max\{l \geq 1 : S[i \dots i+l-1] = S[j \dots j+l-1]\}$. For example, if $S = \text{abracadabra}$, then $lce_S(1, 2) = 0$, $lce_S(0, 3) = 1$, and $lce_S(0, 7) = 4$. Given S in advance for preprocessing, build a data structure for S based on the Karp-Rabin fingerprinting, in $O(n \log n)$ time, so that it supports subsequent online queries of the following two types:

- $lce_S(i, j)$: it computes the longest common extension at positions i and j in $O(\log n)$ time.
- $equal_S(i, j, l)$: it checks if $S[i \dots i+l-1] = S[j \dots j+l-1]$ in constant time.

Analyze the cost and the error probability. The space occupied by the data structure can be $O(n \log n)$ but it is possible to use $O(n)$ space. [Note: in this exercise, a onetime preprocessing is performed, and then many online queries are to be answered on the fly.]

SOLUTION

Karp-Rabin hashing on strings (i.e. $str[n]$) for solving the longest common extension problem. We have the following steps:

1. Create a data structure holding the hashes, with hashes of previous characters being held as a prefix. We fix a sufficient big prime p and we use the Karp-Rabin hash (i.e. for a string k and a base b : $h(k) = (k[0]b^{L-1} + k[1]b^{L-2} + \dots + k[L-1]b^0) \bmod p$):

$$H[0] = h(str[0])$$

$$H[1] = H[0]p + h(str[1])$$

$$H[2] = H[1]p + h(str[2]) = h(str[0])p^2 + h(str[1])p^1 + h(str[2])p^0$$

...

$$H[n-1] = H[n-2]p + h(str[n-1]) = h(str[0])p^{n-1} + \dots + h(str[n-2])p^1 + h(str[n-1])p^0$$

Therefore the space used here is just $O(n)$.

2. Firstly we care about equality. To compare equality of a substring of length l at indexes i, j , we need to know the sub-hash of the string. Thus, we take hashes of $H[i]$ and $H[i+l]$ and we calculate the hash of the sub-string between i and l :

$$\begin{aligned} H[Substring_i] &= H[i+l] - H[i] * p^{(l-1)} \\ &= h(str[0])p^{i+l-1} + \dots + h(str[i+l-2])p^1 + h(str[i+l-1])p^0 \\ &\quad - (h(str[0])p^{i-1} + \dots + h(str[i-2])p^1 + h(str[i-1])p^0) * p^{(l-1)} \\ &= h(str[i+l-1])p^0 + h(str[i+l-2])p^1 \dots + h(str[0])p^{i+l-1} \\ &\quad - (h(str[i-1])p^{(l-1)} + h(str[i-2])p^l \dots + h(str[0])p^{i+l-2}) \\ &= h(str[i+l-1])p^0 + h(str[i+l-2])p^1 \dots + h(str[i])p^{i+l-1} \end{aligned}$$

We do the same procedure for the sub-string between j and l (i.e. $H[Substring_j]$), then we can simply compare the two hash. We introduce a base case (a sanity check) in the case $l = 1$ and we have $str[i] \neq str[j]$. The cost of this operation is $O(1)$ since we did just simple operation $(+)$.

We choose a random prime number $p \in [2, \dots, \tau]$ where $\tau > n$. Since the prime number in the interval $[2, \dots, \tau]$ are approximately $\frac{\tau}{\ln(\tau)}$, and we have a collision when $Substring_i = Substring_j$ but $H[Substring_i] \neq H[Substring_j]$, thus when $c = Substring_i - Substring_j \bmod p = 0$. We can conclude that $P_r[error] \leq \frac{\#BAD_PRIME}{\#PRIME} = \frac{n}{\frac{\tau}{\ln(\tau)}}$ because there are at most n distinct prime p that divide c (Chinese Theorem of residual). If we choose $\tau \approx n^{a+1} \ln(n)$ then we have $P_r[error] \leq \frac{1}{n^a}$.

3. Finally to compute the longest common extension $lce_S(i, j)$, we just do a binary search of the index l . The cost to do so is $O(\ln(n))$ since the check of the equality is constant.

SECOND PROPOSED SOLUTION

Consider a string S of length n , we call $S[i]$ the i -th element in S starting from 0, and $S[i, l]$ the substring of S with length l starting at i (i.e. $S[i] \cdot S[i+1] \dots S[i+l-1]$, where \cdot is the string concatenation).

The idea is to build an array R with the same length of the string S such that for each index i we have that $R[i]$ contains the prefix of S of length $i+1$. $R[0]$ contains the first word of S , $R[1]$ contains the concatenation of the first two words of S etc. In general we will have that $R[i] = S[0, i+1]$.

This representation has of course the problem to be too much expensive, but we will see later how to really implement the data structure, first have a look at how it works for checking the equality of two substrings and for finding the lce_S of two indexes.

Given a pair of indexes i and j , and a length l we want to check whether $S[i, l] = S[j, l]$, looking at R . Considering that $\forall i, l. R[i+l] = S[0, i] \cdot S[i+1, l] = R[i] \cdot S[i+1, l]$, we can prove that

$$\forall i, l. S[i, l] = R[i+l-1] - R[i-1]$$

where we denote by $\alpha - \beta$ the string α without the prefix β (i.e. $\alpha\gamma - \alpha = \gamma$).

Now we can simply check whether $s[i, l] = s[j, l]$ by comparing $R[i+l-1] - R[i-1]$ and $R[j+l-1] - R[j-1]$

The problem of finding the longest common extension is quite simple, it is sufficient to use some form of binary search, exploiting the array R . We use a recursive function in order to find the $lce_S(i, j)$, the input are i, j , the string S and its length n . We consider to have a function $is_equal(A, x, y, l)$ such that return true if $A[x, l] = A[y, l]$. The main difference between this algorithm and a binary search is that we will continue calling recursively the function with a string of halved length regardless of the result of the equality check.

```

function LCE( $S, i, j, n$ )
  if  $n = 1$  then
    if  $is\_equal(S, i, j, 1)$  then
      return 1
    else
      return 0
    end if
  end if
   $l = n/2$ 
  if  $is\_equal(S, i, j, l)$  then
    return  $l + LCE(S, i+l, j+l, n-l)$ 
  else
    return LCE( $S, i, j, l$ )
  end if
end function

```

Note that assuming the is_equal function to have constant cost we have that the cost of lce is $O(\log(n))$ since at each iteration only a recursive call is reached and the length of the input is halved each time.

In the real implementation we will use an array H which is the hash fingerprint of R . For each index i $H[i] = h(R[i]) = R[i] \bmod p$, of course we don't really built the array R , then we can simply fill $H[i]$ with $S[0, i+1] \bmod p$. An implicit assumption in this step (and in Karp-Rabin fingerprint in fact) is to see each string as a number, in particular we see S as the representation of an integer number with base b (depending on the number of possible word in the alphabet). In the following we will call S the string and $(S)_b$ the number for which S is the representation in base b . We will also distinguish between $S[i]$ the word in position i of S and $S[i]_b$ the number between 0 and $b-1$ which represents. Note that we have to redefine h as $h(\alpha) = (\alpha)_b \bmod p$.

Then the number $(S)_b$ will be $S[0]_b \times b^{n-1} + S[1]_b \times b^{n-2} \dots + S[n-1]_b \times b^0$. With this representation we have the beautiful feature that $(S[0, i+l])_b = (S[0, i] \cdot S[i, l])_b = (S[0, i])_b \times b^l + (S[i, l])_b$, that we

can exploit to prove $h(S[i, l]) = H[i + l - 1] - H[i - 1] \times b^l \bmod p$.

$$\begin{aligned}
H[i + l - 1] &= h(R[i + l - 1]) \\
&= h(S[0, i + l]) \\
&= (S[0, i + l])_b \bmod p \\
&= (S[0, i])_b \times b^l + (S[i, l])_b \bmod p \\
&= ((S[0, i])_b \bmod p) \times b^l + ((S[i, l])_b \bmod p) \bmod p \\
&= h(S[0, i]) \times b^l + h(S[i, l]) \bmod p \\
&= h(R[i - 1]) \times b^l + h(S[i, l]) \bmod p \\
&= H[i - 1] \times b^l + h(S[i, l]) \bmod p
\end{aligned}$$

This way the cost of compare two arbitrary substring of the same length $S[i, l]$ and $[j, l]$ is the cost of accessing to four elements of H plus a constant number of arithmetic operations, i.e. $O(1)$. Since we have found a good approximation of the procedure *is_equal*, we can compute *lce_S* with the algorithm shown before in $O(\log n)$ time.

The space occupied by the data structure is quasi $O(n)$ where n is the length of S . Actually the size of the input in bit is n times the size of the word (i.e. $n \times \log b$) and the size of the output is n times the size of an element of H (i.e. $n \times \log p$). Note that we can consider the use of R as a particular case in which p is equal to b^{n+1} (the maximum number that a string of length n with alphabet of size b can represent plus one), in this case the modulus in h is redundant and we have a data structure of size $O(n^2)$.

Note that the trivial way to fill H takes $O(n^2)$. I am referring to something like

```

for  $i \in \{0, 1, \dots, n-1\}$  do
   $H[i] = 0$ 
  for  $j \in \{0, 1, \dots, i\}$  do
     $H[i] = H[i] + S[j] \times b^{i-j} \bmod p$ 
  end for
end for

```

in which we simply implement $H[i] = S[0] \times b^i + S[1] \times b^{i-1} + \dots S[i] \times b^0$. In effect the number of operations (products, sums), is given by $\sum_{i=1}^n i = \frac{n \times (n+1)}{2} = O(n^2)$.

It is not difficult to see that we can use the compositional property of the h function in order to achieve a cost of $O(n)$. The code speaks for itself probably, the idea is to compute $H[0]$ as $S[0] \bmod p$, $H[1]$ as $S[1] + S[0] \times b \bmod p$ etc. We use only a constant number of operations for each word in S .

```

 $H[0] = S[0] \bmod p$ 
for  $i \in \{1, \dots, n-1\}$  do
   $H[i] = H[i-1] \times b + S[i] \bmod p$ 
end for

```

We take p as a random prime number $p \in [2, \dots, \tau]$, where $\tau > n$. Let's have a look at the collision cases: we have a collision when $S[i, l] \neq S[j, l]$ but $h(S[i, l]) = h(S[j, l])$. This is the same as

$$\begin{aligned}
H[i + l - 1] - H[i - 1] \times b^l \bmod p &= H[j + l - 1] - H[j - 1] \times b^l \bmod p \\
H[i + l - 1] - H[i - 1] \times b^l - (H[j + l - 1] - H[j - 1] \times b^l) \bmod p &= 0 \\
h(S[i, l]) - h(S[j, l]) \bmod p &= 0 \\
((S[i, l])_b \bmod p) - (S[j, l])_b \bmod p \bmod p &= 0 \\
(S[i, l])_b - (S[j, l])_b \bmod p &= 0 \\
p \text{ divides } (S[i, l])_b - (S[j, l])_b
\end{aligned}$$

We define $c = (S[i, l])_b - (S[j, l])_b$, we want to count how many bad choices we have for p (how many choices for p such that p divides c). Note that as $(S)_b$ also c is a number representable with a string of length n with alphabet of size b ; then $0 \leq c \leq b^n$. Let k be the number of prime number dividing c , then $c = p_1^{i_1} \times p_2^{i_2} \times \dots \times p_k^{i_k}$ and, since p_x is greater or equal than 2 and i_x is greater or equal than 1 for all x , we can say that $c \geq 2^k$ (as seen in class). Finally $2^k \leq c \leq b^n \leq b^{n \times \log b}$, and $k \leq n \times \log b$.

Since the possible choices for a prime number in the interval $[2, \dots, \tau]$ are approximately $\frac{\tau}{\ln(\tau)}$, the probability of error is less or equal to $\frac{\# \text{BAD PRIMES}}{\# \text{PRIMES}} = \frac{n \times \log b}{\frac{\tau}{\ln(\tau)}}$. If we choose $\tau \approx n^{a+1} \ln(n)$ then we have that the probability of error is less or equal than $\frac{\log b}{n^a}$, and the size of H is approximately $n \times \log p \leq n \times \log(n^{a+1} \times \ln(n)) = n \times (a+1) \times \log(n \times \ln(n)) = O(n \times \log(n \times \log(n)))$ (does it make sense?).

Note that if we want the size of H to be $O(n)$ it is necessary for τ to doesn't depend on n , and in this particular case the probability of error grows up with n .

SOLUTION in $O(N \log N)$ works for every kind of hash function

The proposed data structure that maintain a series of trees. At each level we encode power of two elements (THIS IS IMPOSSIBLE TO EXPLAIN... LOOK THE CODE). The space occupied by this data structure is $O(n \log(n))$ where n is the length of the input array.

```

function CREATETREE( $S, n, p$ )
   $A \leftarrow \text{NEW Array}[\log_2(n) + 1]$ 
   $TEMP \leftarrow \text{NEW Array}[n]$ 
  for  $i$  IN  $(0, n)$  do
     $TEMP[i] \leftarrow S[i] \bmod p$ 
  end for
   $A[0] \leftarrow TEMP$ 
  for  $i$  IN  $(1, \log_2(n))$  do
     $TEMP \leftarrow \text{NEW Array}[n - i]$ 
    for  $j$  IN  $(0, n - i)$  do
       $TEMP[j] \leftarrow A[i - 1][j] + A[i - 1][j + 2^{i-1}] \bmod p$ 
    end for
     $A[i] \leftarrow TEMP$ 
  end for
  return  $A$ 
end function

```

Now to have $lce_S(i, j)$ we build the implement the following procedure, where $h = \lfloor \log_2(n - j) \rfloor$

```

function LCE( $i, j, h$ )
  if  $A[h][i] == A[h][j]$  then
    return  $2^h$ 
  else
    if  $h \neq 0$  then
      return 0
    else
      return  $LCE(i, j, h - 1) + LCE(i + 2^{h-1}, j + 2^{h-1}, h - 1)$ 
    end if
  end if
end function

```

Notice that, this procedure work in $O(\log(n))$ since the array is length is at most $\log(n)$ and we are doing two recursive call with an array one unit smaller each time. Finally to obtain $equalS(i, j, l)$ in cost $O(1)$ we simply check whether $A[\lceil \log_2(l) \rceil][i] == A[\lceil \log_2(l) \rceil][j]$ is true.

Notice that all this algorithm work for arrays in which their length n is a power of two. If we have an array that is not of the latter length, we are doing the following: create the same data structure as before but the part of the array that is not in the tree, that it's at most long $2^{i+1} - 2^i - 1$ where $i = \lfloor \log_2(n) \rfloor$, is store in simple array. In this case whether we need to check if $A[\lfloor \log_2(n - j) \rfloor][i] == A[\lfloor \log_2(n - j) \rfloor][j]$, i.e. the longest possible string, we need also to check, manually, whether there is a matching in the array. Last but not least, the calculation of the error probability is exactly the same to the one analysed during the course.