

1-D range query

Describe how to efficiently perform one-dimensional range queries for the data structures described in Problems 14 and 15. Given two keys $k_1 \leq k_2$, a range query asks to report all the keys k such that $k_1 \leq k \leq k_2$. Give an analysis of the cost of the proposed algorithm, asking yourself whether it is output-sensitive, namely, it takes $O(\log_B N + R/B)$ block transfers where R is the number of reported keys.

SOLUTION

The main idea in this exercise is to exploit the fact that keys at each level are stored sequentially, and they are sorted since we have a binary search tree¹. Therefore, we can query the two keys k_1, k_2 using a simple binary search, and take all the keys that are in this range. Let's go in details for both data structures.

Problem 14 In this case the tree has been stored in an array, in which every layer of the tree is stored sequentially. For example for a binary tree we have:

$$[root, L_0, R_1, L_{00}, R_{01}, L_{10}, R_{11}, \dots]$$

Therefore, the algorithm is going to search simultaneously the two keys k_1, k_2 , and thus at each level is going through each layer. Here is going to find two position t_1, t_2 (in the case are not directly the keys) in which the algorithm is going to land, then all the keys we will find in the middle of t_1, t_2 are in the range k_1, k_2 .

Problem 15 Here is a bit less intuitive but we exploit the same fact. Ones we have store implicitly the tree we are going to have a recursive layout of the tree, as shown in the picture below.




FIG. 2.1. The van Emde Boas layout. Left: in general; right: of a tree of height 5.

Therefore at each branching of the b-search, we are going to land (in the worst case) in two different blocks, and then we get all the keys that we encounter while looking for k_1, k_2 until we reach the keys in themselves. At this point we get all the blocks that are in the middle.

Notice that at certain point, in both cases we are going to end up with a sequence of key (R) stored sequentially. Therefore, we are going to have $2\log_B N$ access for two binary search and another $O(\frac{R}{B})$ to return the keys. This means that the algorithm is output-sensitive.

¹Notice that in both exercise we have provided a way to access the left and the right child