

External memory mergesort

In the external-memory model (hereafter EM model), show how to implement the k -way merge (where $(k+1)B \leq M$), namely, how to simultaneously merge k sorted sequences of total length N , with an I/O cost of $O(N/B)$ where B is the block transfer size. Also, try to minimize and analyze the CPU time cost.

SOLUTION

The classical external sorting algorithm, called merge-sort, consists of two phases:

1. **The sort phase.** B element are read into the buffer, sorted, and written to the disk. This creates $n = \lceil N/B \rceil$ sorted subset of records, called runs, stored in separate auxiliary files, numbered from 1 to n . The runs have all the same number of pages, B , except the last.
2. **The merge phase** consists of multiple merge passes. In each merge pass, $k = B - 1$ runs are merged using the remaining buffer page for output. At the end of a merge pass, the number of runs becomes $n = \lceil n/k \rceil$. A merge pass is repeated until $n > 1$. The merge is done transferring k blocks in main memory from the n , keeping the minimum from each run, and store it in another block (notice is the reason for $k+1$ block). When we full fill the latter block we transfer it in external memory.

This is possible because $(k+1)B \leq M$.




Figure 1: Let us show how to sort the file A_0 containing 12 pages, with file and buffer pages capacity of 2 records[IT'S AN EXAMPLE no need to consider the records], $B = 3$ and 2-merge passes.

- The initial sort phase creates the runs A_1, A_2, A_3 and A_4 .
- The first merge pass creates the runs A_5 and A_6 by merging A_1, A_2 and A_3, A_4 .
- The second merge pass creates the sorted data A_7 by merging A_5, A_6 .

At each iteration of merge we read $\lceil N/B \rceil$ block and the same for the writing. Therefore we have $O(N/B)$ I/O operation.

The bottle neck in CPU cost is to search of the min key in the k runs each time we add an element in the output block. Indeed if we do a linear search we pay $O(k)$ at each element, then a total of $O(Nk)$. Instead, if we use an MinHeap (priority queue), we pay $O(\log k)$ to insert an element in the Heap and $O(1)$ to retrieve the min. Therefore the total cost in $O(N \log k)$.